

Arjun Chavan

Tysons Corner, VA · arjun.chavan999@gmail.com · +1 650-338-9658 · [LinkedIn](#) · [GitHub](#) · [Portfolio](#)

SUMMARY

Software Engineer with 3+ years building distributed backend systems and cloud infrastructure at scale — including the LLM-serving layer of an enterprise AI product line handling 2M+ daily queries and REST platforms serving 5M+ daily requests. Deep Kubernetes/AWS/Terraform operations background with a track record of measurable performance wins (90% deploy-time reduction, 60% DB-load reduction, ~\$50K/yr infra savings). Recently promoted to SWE2 at Strategy; scope includes production ownership, customer-facing deployments, and mentoring 4 engineers.

WORK EXPERIENCE

Strategy — Tysons Corner, VA

Software Engineer 2 · Apr 2026 – Present

Software Engineer 1 · Feb 2025 – Apr 2026

Software Engineering Intern · May 2024 – Aug 2024

- Built the **application-layer LLM orchestration infrastructure** for production AI products — multi-provider integration (OpenAI, Groq, Anthropic) sustaining **2M+ daily queries** with provider failover
- Designed a **dynamic multi-provider routing layer** cutting AI infrastructure cost by **~\$50K/year** — latency-, cost-, and availability-aware dispatch across providers
- Led design of **Autodash 2.0** end-to-end — PoC through production launch of a multi-agent workflow system for autonomous dashboard creation
- Drove systematic **prompt/token optimization** across enterprise AI applications: **40% token reduction with 92% accuracy retained** (eval-verified)
- (Post-promotion)** Own customer-facing deployment engineering for enterprise AI features — **Kubernetes on GCP and Docker-based Azure deployments** for major enterprise clients
- (Post-promotion)** Mentoring 4 incoming engineers (2 interns, 2 new-grad hires); recognized at Strategy's 2026 Annual Conference for shipping a functional AI use case in under 3 days

Cimpress — Mumbai, India

DevOps Engineer · Jan 2022 – Aug 2023

Software Engineer (Full Stack) · Aug 2021 – Dec 2021

- Managed **15+ AWS services** including **EKS** and CodePipeline supporting production deployments serving **500K users**
- Automated CI/CD with Docker and Kubernetes, cutting deployment time by **90%** with **zero-downtime releases**
- Introduced **Terraform IaC across 4 teams**, achieving 100% deployment consistency and eliminating configuration drift
- Owned 5 production infrastructure projects end-to-end, including the team's first **Docker/K8s CronJob** service (CMS backup verification) and a Git/CodeCommit-internals rollback strategy
- Key role in the **NA Migration** company-wide launch — time-sensitive multi-team deployment; recognized for execution under pressure
- (Full Stack)** Built **REST APIs handling 5M+ daily requests**; introduced caching that **reduced DB load by 60%**; raised code coverage 45% → 82%

PROJECTS

Market Sentiment RAG — Grounded Retrieval Pipeline · *personal* · github.com/ArjunChavan219/market-sentiment-rag

- Qwen3** over **Qdrant**; **every claim cites its source passage**, enforced by a per-request JSON schema — removing it dropped the model to **zero** valid citations. **FinBERT**-baselined evals + scorecard; **vLLM/Vertex** code written, **not deployed**

LibreChat CourseTools — AI Course Management Platform · *grad project, JHU*

- Architected role-based AI chat platform serving 1,000+ students; multi-model backend with professor/TA/student auth hierarchies

EDUCATION

MS Computer Science · Johns Hopkins University, MD · GPA 3.9/4.0 · Dec 2024

- Coursework: Software System Design, Cloud Computing, Neural Networks, AI, Deep Learning · Course Assistant: AI (×2), Deep Learning

BE Computer Engineering · Mumbai University, India · GPA 3.9/4.0 · Jun 2021 · Best Outgoing Student

SKILLS

Languages: Python, C/C++, TypeScript / JavaScript (Node.js), Java, SQL

Distributed / Infra: Kubernetes, Docker, Terraform, AWS (EKS, CodePipeline, S3, CloudWatch), GCP (Vertex AI), RabbitMQ, Celery, async pipelines, CI/CD, zero-downtime deploys

AI Serving (applied): self-hosted model serving (vLLM, Vertex AI), multi-provider LLM orchestration, OpenAI-compatible APIs, RAG, vector databases (Qdrant), cost/latency optimization, eval design

Databases: PostgreSQL, MySQL, MongoDB, DynamoDB, ClickHouse

Practices: system design, observability & on-call ownership, IaC, customer-facing deployment engineering, mentorship